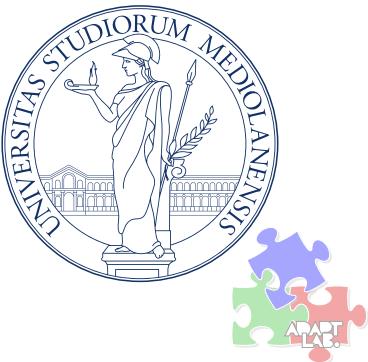
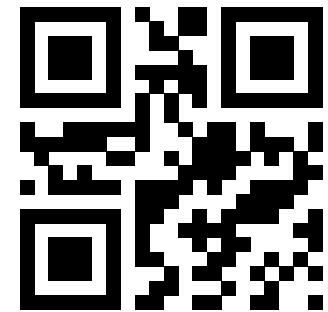


MLIR: Scaling Compiler Infrastructure for Domain Specific Computation [1]



Federico Bruzzone,¹ PhD Candidate

Milan, Italy – 18 March 2026



¹ADAPT Lab – University of Milan,
Website: federicobruzzone.github.io,
Github: github.com/FedericoBruzzone,
Email: federico.bruzzone@unimi.it
Slides: [TODO](#)

MLIR: Multi-Level Intermediate Representation

Part of the LLVM project, the MLIR is a novel approach to building **reusable**, **modular**, and **extensible** compiler infrastructure.

MLIR aims to address software fragmentation, improve compilation for heterogeneous hardware, significantly reduce the cost of building **domain specific compilers**, and aid in connecting existing compilers together.





Why another compiler infrastructure?

Although the *one size fits all* approach of traditional compilers (e.g., LLVM [2] and JVM [3]) has been successful for general-purpose programming, it has shown limitations in the context of domain-specific applications.

Many problems are better modeled at a **higher-** or **lower-level abstraction** – e.g., source-level static analysis of C++/Rust is difficult on LLVM IR.

Hence, many languages and frameworks developed their own intermediate representations (IRs) to leverage the **semantic information** of their domain – including TensorFlow’s XLA HLO, PyTorch’s Glow, Rust’s MIR, Swift’s SIL, Clang’s CIL, and so on.

While domain-specific IRs are well-understood, their *high engineering costs* often lead to compromised infrastructure quality. This results in *suboptimal compilers* plagued by bugs, latency, and a poor debugging experience [1].



MLIR to the rescue

MLIR directly addresses these issues by making it **cheap** to design and **introduce** new abstraction layers.

It achieves this by:

- standardizing the Static Single Assignment (SSA)-based IR data structures,
- providing a declarative system for defining IR *dialects*, and
- providing a wide range of common infrastructure including documentation, parsing and printing logic, location tracking, multithreaded compilation support, pass management.





Design Principles

- **Parsimony**: Apply *Occam's razor* to builtin semantics, concepts, and programming interface. Specify invariants once, but verify correctness throughout \Rightarrow *extensibility*.
- **Traceability**: Retain rather than recover information. Declare rules and properties to enable transformation, rather than step wise imperative specification \Rightarrow *composability*.
- **Progressivity**: Premature lowering is the root of all evil. Beyond representation layers, allow multiple transformation paths that lower individual regions on demand \Rightarrow *reusability*.



Little Built-in, Everything Customizable [Parsimony]

- The system is based on a minimal number of fundamental concepts, leaving most of the intermediate representation fully **customizable**.
- A handful of abstractions—types, operations and attributes—should be used to express *everything else*, allowing fewer and more consistent abstractions that are easy to **comprehend**, **extend**, and **adopt**.
- A success criterion for customization is the possibility to express a diverse set of abstractions including **ML graphs**, ASTs, mathematical abstractions such as **polyhedral**, CFGs and instruction-level IRs such as **LLVM IR**, without hard-coding concepts.



SSA and Regions [*Parsimony*]

- SSA [4] makes dataflow analysis *simple* and *sparse*. However, while many existing IRs use this flat, linearized CFG, representing higher level abstractions push introducing **nested regions**² as a first-class citizen – e.g., structured control flow, concurrency constructs, and closures.
- The (LLVM) normalization/canonicalization process is sacrificed due to the presence of multiple ways to represent the same semantics.
- The frontend is responsible for choosing the level of abstraction for the IR.

²A region is a single-entry, multi-exit CFG that can be nested inside an operation. It is a generalization of the concept of basic blocks and allows for more flexible control flow representation.



The Canonical Loop Structure

Pre-header, header, latch, and body is a prototypical loop structure.

```
; for (int i = 0; i < n; ++i) { ... }

entry:
    br label %header

header:
    %i = phi i32 [ 0, %entry ], [ %inc, %latch ]
    %cmp = icmp slt i32 %i, %n
    br i1 %cmp, label %body, label %multi-exit

latch:
    %inc = add i32 %i, 1
    br label %header

body:
    ; loop body
    br label %latch

multi-exit:
    ; code after the loop
```

llvm

```
// A simple loop from 0 to 10 with a step of 1
scf.for %i = %c0 to %c10 step %c1 {
    // Loop body goes here
    // %i is the induction variable
    "some.operation"(%i) : (index) -> ()
}
```

mlir

```
// An affine loop: optimized for polyhedral compilation
affine.for %i = 0 to 10 {
    %val = affine.load %buffer[%i] : memref<10xf32>
    // ... operations ...
}
```

mlir



Maintain Higher-Level Semantics [*Progressivity*]

- Attempts to **recover** abstract semantics once lowered are fragile and often **fail** to capture the full semantics.
- The system should maintain the structure of computations and **progressively lower** to the hardware abstraction.
- Removing structured control flow – i.e. lowering to a CFG – essentially means no further transformations will be performed that exploits the structure.
- Previous compilers have been introducing multiple fixed levels of abstraction in their pipeline causing **phase ordering** issues.



Declaration and Validation [*Parsimony/Traceability*]

- Defining representation modifiers should be as simple as introducing new abstractions.
- Common transformations should be implementable as **rewrite rules** expressed declaratively.
- Although rewriting systems are well-studied, the MLIR's extensibility opens up new challenges.
- While verification, testing, and translation validation [5] are useful a more robust approach to combining all these techniques for **extensible** and **modular** IRs.





Source Location Tracking [*Traceability*]

- **Lack-of-transparency** in complex compilation systems is ubiquitous. This is particularly problematic when compiling safety-critical and sensitive applications (cf. WYSINWYX by Balakrishnan et al. [6]).
- Thus, the **provenance** of an operation — including its original location and applied transformations — should be easily traceable within MLIR.
- One indirect goal of accurately propagating high-level information to the lower levels is to help support **secure** and **traceable** compilation.

Intermediate Representations Design

MLIR has a *generic* textual representation that supports MLIR's extensibility and fully reflects the in-memory representation, which is paramount for traceability, manual IR validation and testing. Extensibility comes with

```
// Attribute aliases can be forward-declared.  
#map1 = (d0, d1) -> (d0 + d1)  
#map3 = ()[s0] -> (s0)  
  
// Ops may have regions attached.  
"affine.for"(%arg0) {  
// Regions consist of a CFG of blocks with arguments.  
^bb0(%arg4: index):  
    // Block are lists of operations.  
    "affine.for"(%arg0) {  
        ^bb0(%arg5: index):  
            // Ops use and define typed values, which obey SSA.  
            %0 = "affine.load"(%arg1, %arg4) {map = (d0) -> (d0)}  
                : (memref<?xf32>, index) -> f32  
            %1 = "affine.load"(%arg2, %arg5) {map = (d0) -> (d0)}  
                : (memref<?xf32>, index) -> f32  
            %2 = "std.mulf"(%, %1) : (f32, f32) -> f32  
            %3 = "affine.load"(%arg3, %arg4, %arg5) {map = #map1}  
                : (memref<?xf32>, index, index) -> f32  
            %4 = "std.addf"(%, %2) : (f32, f32) -> f32  
            "affine.store"(%4, %arg3, %arg4, %arg5) {map = #map1}  
                : (f32, memref<?xf32>, index, index) -> ()  
        // Blocks end with a terminator Op.  
        "affine.terminator"(): () -> ()  
    // Ops have a list of attributes.  
} } {lower_bound = () -> (0), step = 1 : index, upper_bound = #map3}  
    : (index) -> ()  
    "affine.terminator"(): () -> ()  
} } {lower_bound = () -> (0), step = 1 : index, upper_bound = #map3}  
    : (index) -> ()
```

the burden of verbosity,
which can be
compensated by the
custom syntax that MLIR
supports

Thank You!



Bibliography

- [1] C. Lattner *et al.*, “MLIR: Scaling Compiler Infrastructure for Domain Specific Computation,” in *2021 IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*, 2021, pp. 2–14. doi: [10.1109/CGO51591.2021.9370308](https://doi.org/10.1109/CGO51591.2021.9370308).
- [2] C. Lattner and V. Adve, “LLVM: A compilation framework for lifelong program analysis & transformation,” in *International symposium on code generation and optimization, 2004. CGO 2004.*, 2004, pp. 75–86.
- [3] T. Lindholm, F. Yellin, G. Bracha, and A. Buckley, *The Java virtual machine specification*. Addison-wesley, 2013.
- [4] R. Cytron, J. Ferrante, B. K. Rosen, M. N. Wegman, and F. K. Zadeck, “Efficiently computing static single assignment form and the control dependence graph,” *ACM Trans. Program. Lang. Syst.*, vol. 13, no. 4, pp. 451–490, Oct. 1991, doi: [10.1145/115372.115320](https://doi.org/10.1145/115372.115320).
- [5] A. Pnueli, M. Siegel, and E. Singerman, “Translation validation,” in *Tools and Algorithms for the Construction and Analysis of Systems*, B. Steffen, Ed., Berlin, Heidelberg: Springer Berlin Heidelberg, 1998, pp. 151–166.
- [6] G. Balakrishnan and T. Reps, “WYSINWYX: What you see is not what you eXecute,” *ACM Trans. Program. Lang. Syst.*, vol. 32, no. 6, Aug. 2010, doi: [10.1145/1749608.1749612](https://doi.org/10.1145/1749608.1749612).